

# CS111 Winter 2024

---

Discussion 1C/F (week4)

Yadi Cao

# Logistics

- Lab 1 is due this Fri (Apr/27)
- Last Dis: **Hints** for Lab1
- This Dis: Lab 2 background knowledge
- Next Fri Dis (Apr/03): **Lab 02 Code skeleton**

# Agenda

- What is scheduling, why do we need a good scheduling?
- Typical scheduling policies, PROs and CONs.
- Lab2 background (Round-robin, concrete example)
- Lab2 starter (coding)

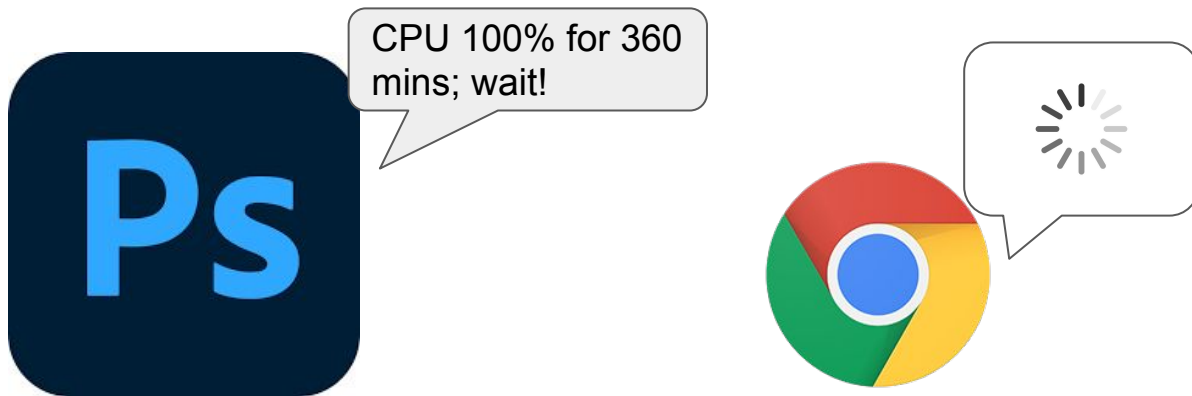
# Lab2 Background

---

# Why do we need a scheduling?

Can you depict a modern computer without scheduling?

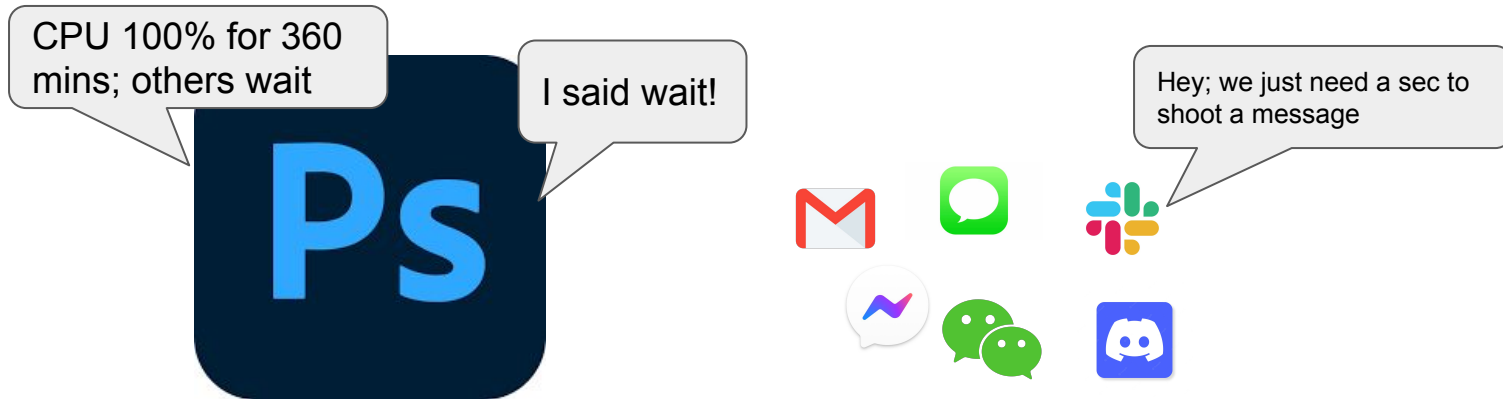
- You are opening a browser to get some picture resources, but the background PhotoShop is rendering forever...you have to wait for PS to finish



- Formal ans: because CPU is a **preemptable resources**; sharing CPU time between processes can increase **efficiency and fairness**; crucial for **concurrency**

# Why do we need a good scheduling?

- The scheduler obeys the strict ordering of incoming tasks



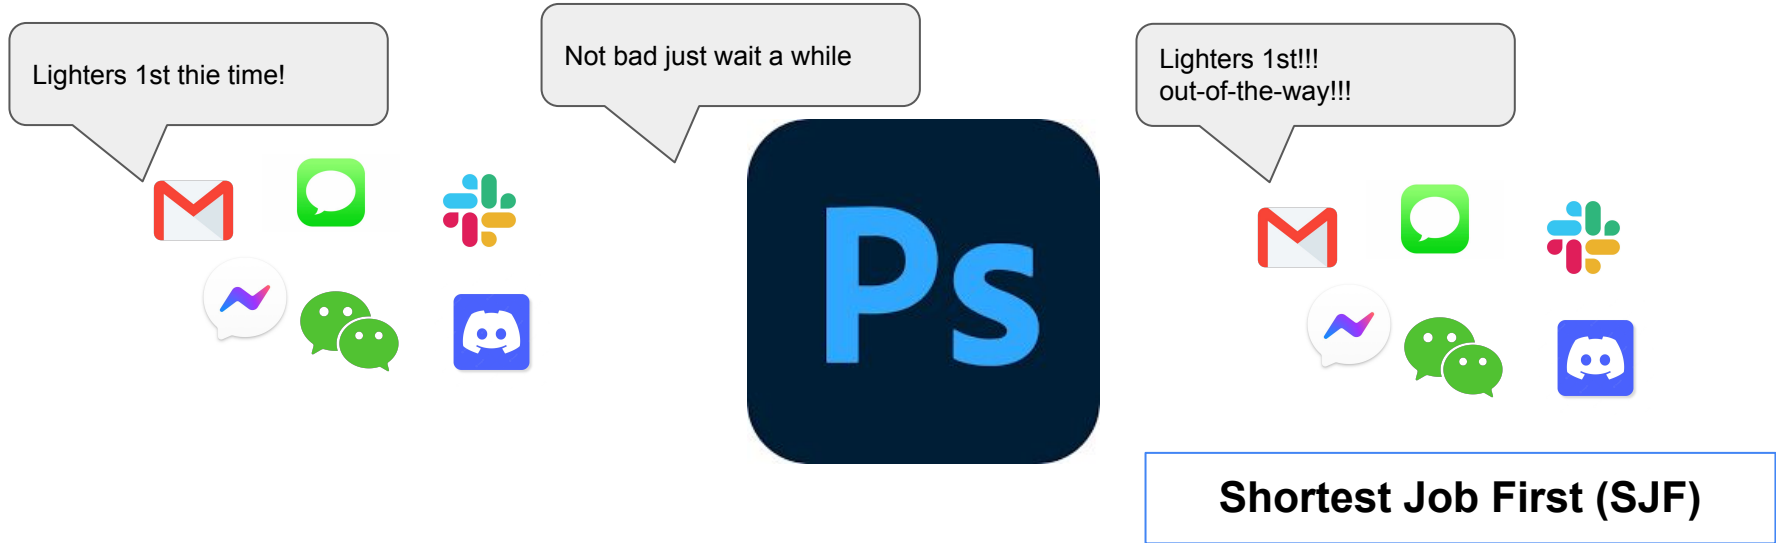
Problems?

**First Come First Served  
(FCFS)**

- so even tiny tasks like email takes 0.5s to send, if they arrive late, they have to wait
- Formal ans: **wait and response** time will be too long

# Why do we need a good scheduling?

- The scheduler now is smarter, it always picks up the easiest job 1st



Problems?

- The heavy job never gets even started
- Formal ans: **starves** heavy jobs

# Shortest Job First (SJF), continued

- Always schedule the job with the shortest burst time first
- No-preemption (don't interrupt the running process)

Provable, this is optimal at minimizing average wait time

Besides **starving**, any other practical problems?

- CPU doesn't know the burst time before running a task (estimate)



# Shortest Remaining Time First (SRTF)

- Improved SJF with preemption:
  - When a new task added, compare its burst time with current task remaining time.
  - If the new task has smaller burst time, swap.

Same problems as SJF!

- CPU doesn't know the remaining time until a task finishes (then remaining time is zero)
- Starving the heavy jobs

# Metrics

- Minimize waiting time and response time  
Don't have a process waiting too long (or too long to start)
- Maximize CPU utilization  
Don't have the CPU idle
- Maximize throughput  
Complete as many processes as possible
- Fairness  
Try to give each process the same percentage of the CPU

# What is Lab 2 about?

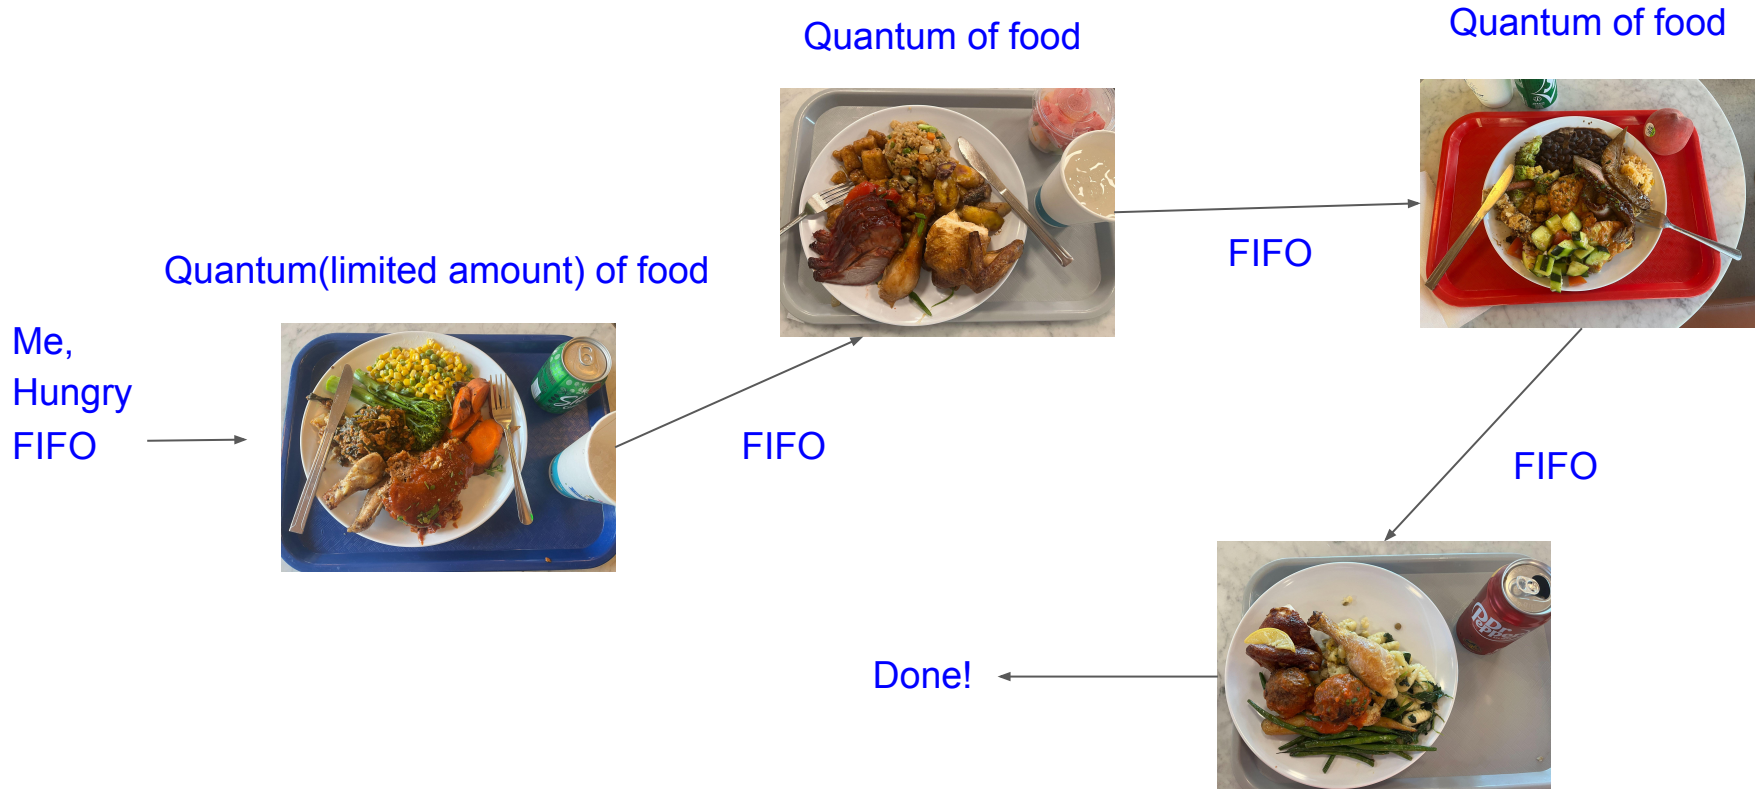
- **Implement round-robin (RR) scheduling**

- You need to write “**a simulation**” for **round-robin (RR) scheduling** for the specified workload and the **quantum length**.
- You’ll be given **a basic skeleton** that parses an input file and command line arguments.
- To be able to do this lab successfully, you’ll need to understand how round-robin scheduling works and utilize a data structures(**which?**) needed.
- The output is **average waiting and response time** for this “**simulation run**”

- **Example output**

```
cs111@cs111 cs111/lab2-pan (main *) » make
cc -std=gnu17 -Wpedantic -Wall -O0 -pipe -fno-plt -fPIC -c -o rr.o rr.c
cc -lrt -Wl,-O1,--sort-common,--as-needed,-z,relro,-z,now rr.o -o rr
cs111@cs111 cs111/lab2-pan (main *) » ./rr processes.txt 3
Average waiting time: 7.00
Average response time: 2.75
```

# UCLA dining hall is Round Robin



# Round Robin

- **Round Robin** is a scheduling algorithm that optimizes **fairness** and **response time**
  - The operating system divides execution into time slices (or quanta), and an individual time slice is called a **quantum**
  - It maintains a **FIFO** (First in First Out) **queue** of processes
  - A process will be **preempted** if it is still running at end of quantum and **re-add to the end of the queue**
- **Round Robin** is similar to **Dining hall in Google**
  - Everyone first comes to the hall in a **FIFO queue**
  - The staff will let you sequentially pick up a **quantum types of food (also limited amount)**
  - At the end of the **quantum**, either you are **done**, satisfied with the food and **leave**
  - If not, say you need more shrimp, **rejoin the end of queue** (also hide your plate), and pick up a second round

# Round Robin Example (in exam 2 years ago)

Assume RR with a quantum length of 3 units, and the 4 tasks arrive/run(burst) time are listed as below

Process	Arrival Time	Burst Time
P <sub>1</sub>	0	7
P <sub>2</sub>	2	4
P <sub>3</sub>	4	1
P <sub>4</sub>	5	4

- Q1: Draw the trajectory of occupancy(which task) of CPU
- Q2: Draw the trajectory of the Queue of the unfinished tasks
- Q3: Calculate the average wait/response time

# Definitions

`waiting_time = end_time - arrival_time - burst_time`

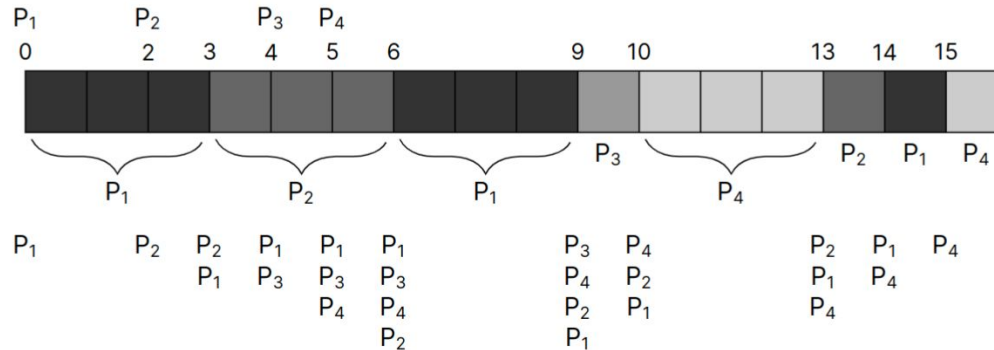
`response_time = start_execute_time - arrival_time`

# Draw animation and replace the below

RR with a quantum length of 3 units

Process	Arrival Time	Burst Time
P <sub>1</sub>	0	7
P <sub>2</sub>	2	4
P <sub>3</sub>	4	1
P <sub>4</sub>	5	4

For RR, our schedule is (arrival on top, queue on bottom):

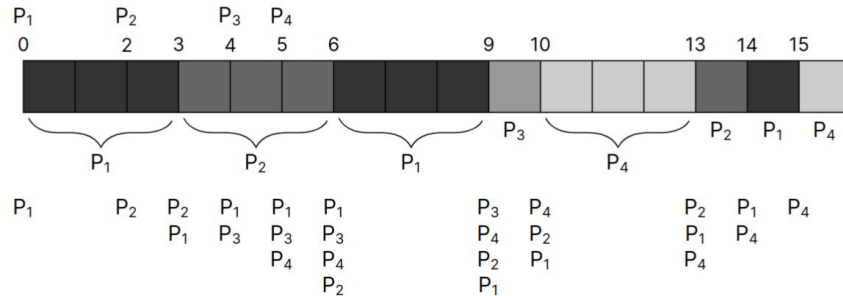




# Round Robin Example

Process	Arrival Time	Burst Time
P <sub>1</sub>	0	7
P <sub>2</sub>	2	4
P <sub>3</sub>	4	1
P <sub>4</sub>	5	4

For RR, our schedule is (arrival on top, queue on bottom):



Note: processes with same arrival time should be added to queue in order

T=0, P<sub>1</sub> arrives, add P<sub>1</sub> the front of the queue

T=2, P<sub>2</sub> comes in, added to the queue

T=3, P<sub>1</sub> is done, P<sub>2</sub> is schedule 3 units, P<sub>1</sub> is on top now

T=4, P<sub>3</sub> comes in, sent to the back of the queue

T=5, P<sub>4</sub> comes in, sent to the back of the queue

T=6, P<sub>2</sub> is done, added to bottom, P<sub>1</sub> is on top and runs

T=9, P<sub>1</sub> is done, P<sub>3</sub> is on top and runs

T=10, P<sub>3</sub> is done, P<sub>4</sub> is on top and runs

T=13, P<sub>4</sub> goes to the back, P<sub>2</sub> is on top and runs

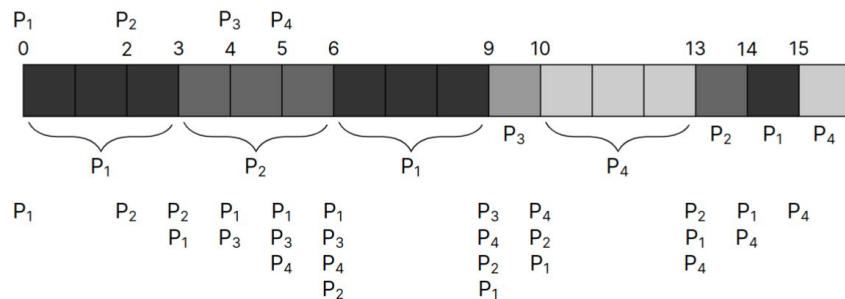
T=14, P<sub>1</sub> is on top and runs with 1 unit

T=15, P<sub>4</sub> is left and runs with 1 unit

# Round Robin Example

Process	Arrival Time	Burst Time
P <sub>1</sub>	0	7
P <sub>2</sub>	2	4
P <sub>3</sub>	4	1
P <sub>4</sub>	5	4

For RR, our schedule is (arrival on top, queue on bottom):



## Average waiting time

The elapsed time a process is waiting for other processes since its arrival

P<sub>1</sub>:

- finishes at 15, starts at 0 and it takes 7
- $15 - 0 - 7 = 8$

P<sub>2</sub>:

- finishes at 14, starts at 2 and it takes 4
- $14 - 2 - 4 = 8$

P<sub>3</sub>:

- finishes at 10, starts at 4 and it takes 1
- $10 - 4 - 1 = 5$

P<sub>4</sub>:

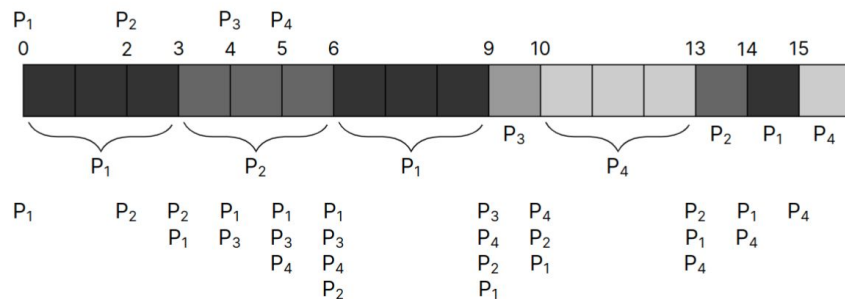
- finishes at 16, starts at 5 and it takes 4
- $16 - 5 - 4 = 7$

$$\text{Average waiting time: } \frac{8+8+5+7}{4} = 7$$

# Round Robin Example

Process	Arrival Time	Burst Time
P <sub>1</sub>	0	7
P <sub>2</sub>	2	4
P <sub>3</sub>	4	1
P <sub>4</sub>	5	4

For RR, our schedule is (arrival on top, queue on bottom):



## Average response time

The time starting from its first arrival to when it first starts executing

P1:

- It starts at 0, arrives at 0
- 0

P2:

- It starts 3, arrives at 2
- $3-2=1$

P3:

- It starts at 9, arrives at 4
- $9-4=5$

P4:

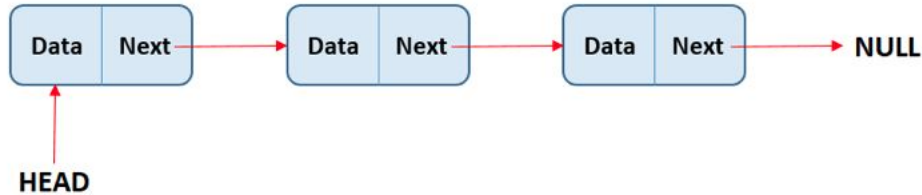
- It starts at 10, arrives at 5
- $10-5=5$

$$\text{Average response time: } \frac{0+1+5+5}{4} = 2.75$$

# Now we want to implement, which data structure?

## FIFO queue by linked list

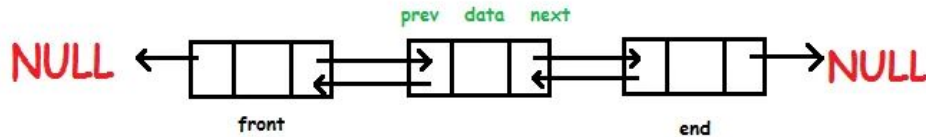
- A linked list is a **linear dynamic** data structure to store data items.
- Unlike arrays, the linked list does not store data items in contiguous memory locations
- Instead, the elements in a linked list are “**linked**” together using **pointers**.



# Doubly Linked List

## What is a doubly linked list (DLL)

- A type of linked list in which each node contains data as well as two pointers.
- One to the next, another to the previous container
- The head node's previous ptr points to NULL
- The tail node's next ptr points to NULL



The two pointers help us to **traverse the list in both backward and forward direction.**

# Great but we don't have linked list in C

We are writing kernel code, we use C.

**C doesn't have a linked list type, unlike C++.** What are your options ?

## Option 1: Implement it ourselves

- We might need to define a node struct (i.e, the element), then lookup, delete, insert, length methods to access them
- Of course, this requires a lot of “fun” pointer operations

## Option 2: “steal” something already existed

- We just need to import it and use it !
- But in order to use it correctly, we need to understand it correctly!

**tailq**

See the appendix :->



*Don't Reinvent The Wheel, Unless You Plan on Learning More About Wheels*

# Lab 2 Skeleton

---

# Lab02 Overview

- Implement part of a scheduling “**simulation**”
  - Schedule processes based on Round-Robin algorithm
  - Implement scheduling queue using doubly-linked list (TAILQ macros)
  - Calculate average waiting and response times
- Example Output

```
cs111@cs111 cs111/lab2-pan (main *) » make
cc -std=gnu17 -Wpedantic -Wall -O0 -pipe -fno-plt -fPIC -c -o rr.o rr.c
cc -lrt -Wl,-O1,--sort-common,--as-needed,-z,relro,-z,now rr.o -o rr
cs111@cs111 cs111/lab2-pan (main *) » ./rr processes.txt 3
Average waiting time: 7.00
Average response time: 2.75
```



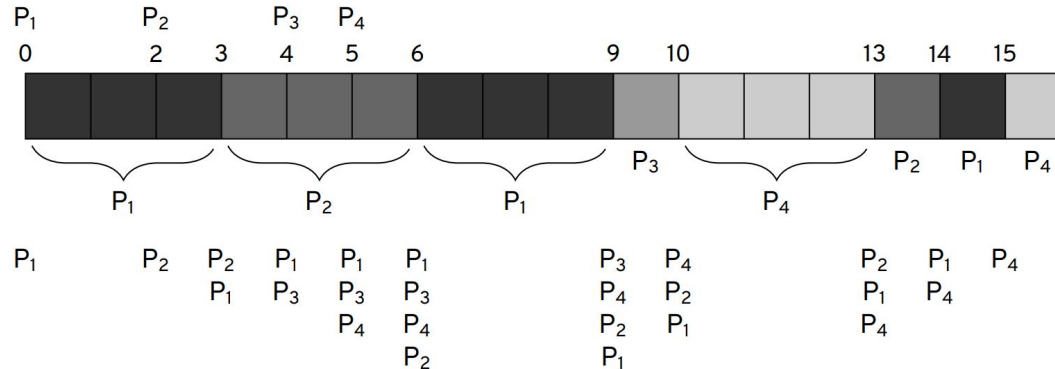
# Process.txt

```
4
1, 0, 7
2, 2, 4
3, 4, 1
4, 5, 4
```



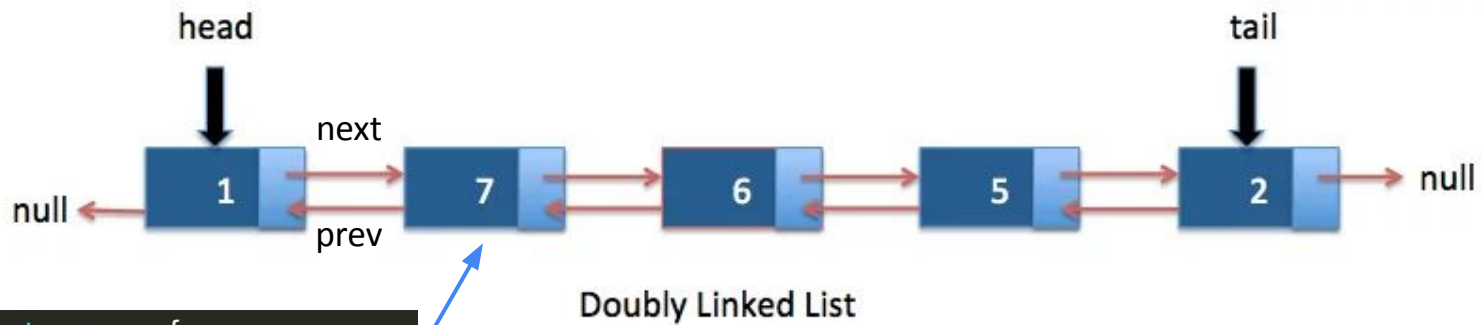
Process	Arrival Time	Burst Time
P <sub>1</sub>	0	7
P <sub>2</sub>	2	4
P <sub>3</sub>	4	1
P <sub>4</sub>	5	4

For RR, our schedule is (arrival on top, queue on bottom):



# Doubly Linked List

**struct process {}** declares a process node structure on a doubly linked list (queue)



```
struct process {  
    u32 pid;  
    u32 arrival_time;  
    u32 burst_time;  
  
    u32 remaining_time;  
    u32 start_exec_time;  
    u32 waiting_time;  
    u32 response_time;  
  
    TAILQ_ENTRY(process) pointers;  
};
```

# Skeleton code

```
1  #include <errno.h>
2  #include <fcntl.h>
3  #include <stdbool.h>
4  #include <stdint.h>
5  #include <stdio.h>
6  #include <stdlib.h>
7  #include <sys/queue.h>
8  #include <sys/mman.h>
9  #include <sys/stat.h>
10 #include <unistd.h>
11
12 typedef uint32_t u32;
13 typedef int32_t i32;
14
15 struct process {
16     u32 pid;
17     u32 arrival_time;
18     u32 burst_time;
19
20     TAILQ_ENTRY(process) pointers;
21
22     /* Additional fields here */
23     /* End of "Additional fields here" */
24 };
25
26 TAILQ_HEAD(process_list, process);
```

```
127 int main(int argc, char *argv[])
128 {
129     if (argc != 3) {
130         return EINVAL;
131     }
132     struct process *data;
133     u32 size;
134     init_processes(argv[1], &data, &size);
135
136     u32 quantum_length = next_int_from_c_str(argv[2]);
137
138     struct process_list list;
139     TAILQ_INIT(&list);
140
141     u32 total_waiting_time = 0;
142     u32 total_response_time = 0;
143
144     /* Your code here */
145     /* End of "Your code here" */
146
147     printf("Average waiting time: %.2f\n", (float) total_waiting_time / (float) size);
148     printf("Average response time: %.2f\n", (float) total_response_time / (float) size);
149
150     free(data);
151     return 0;
152 }
153
```

`&data` is the array of all processes parsed from txt  
`size` is the number of processes from txt

```
init_processes(argv[1], &data, &size);
```



`./rr processes.txt 3`

`processes.txt`

*Arrival time will  
not always be  
in order*

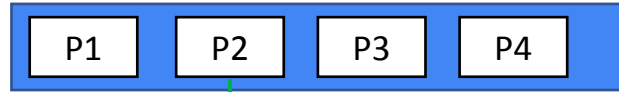
`quantum_time`

*Should be >0  
Invalid arg exit with EINVAL (optional)*

`While(!finished)`

```
struct process *data;
```

All processes parsed from txt



Check arrival time (t)  
Add newly arrived process to queue

*If multiple process, insert based  
on top to bottom order in txt*

```
struct process_list list;
```

`list` is the scheduling queue  
of active processes

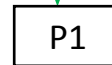


*Queue can be  
empty*

*After newly  
arrived  
process*

pop the 1st process  
in the queue

```
Remove process from head of queue  
(pop may have done the removal)  
If (remaining time > 0){  
    insert unfinished process to end of queue  
}
```



Calculate avg  
wait time &  
response time

```
If all  
data[i].remaining  
time == 0  
Break while; finished
```

- Run process for 1 unit at a time for a total of  $\min(\text{quantum length, remaining time})$
- Decrease process remaining time
- Check loop termination condition
- `++t`

# Creating the queue

**TAILQ\_ENTRY**(NODE);

- Define **entries to the next/ previous NODE** in the queue
- **NODE** is the user defined node structure

**TAILQ\_HEAD**(CONTAINER, NODE);

- Define a **container** that stores **all** the **NODEs** in the queue
- **CONTAINER** is a newly defined **structure** (e.g. use `struct CONTAINER node_list`; to initialize a **CONTAINER** instance `node_list`)
- Always followed by **TAILQ\_INIT**(&node\_list); this when the initialization of the queue occurs

```
struct process {  
    u32 pid;  
    u32 arrival_time;  
    u32 burst_time;  
  
    TAILQ_ENTRY(process) pointers; //macro for creating pointer  
};  
  
TAILQ_HEAD(process_list, process);
```



```
int main(int argc, char *argv[])  
{  
    //...  
  
    //list (queue) of active processes  
    struct process_list list;  
    TAILQ_INIT(&list);  
}
```

# Modifying the queue

```
struct process {
    u32 pid;
    u32 arrival_time;
    u32 burst_time;

    TAILQ_ENTRY(process) pointers; //macro for creating pointer
};

TAILQ_HEAD(process_list, process);
```

```
int main(int argc, char *argv[])
{
    //...

    //list (queue) of active processes
    struct process_list list;
    TAILQ_INIT(&list);
}
```

- void TAILQ\_INSERT\_TAIL(TAILQ\_HEAD \*head, struct NODE \*elm, TAILQ\_ENTRY ENTRY);

```
struct process *p;
p = &data[i];

TAILQ_INSERT_TAIL(&list, p, pointers);
```

- TAILQ\_FOREACH(struct NODE \*var, TAILQ\_HEAD \*head, TAILQ\_ENTRY ENTRY);

```
TAILQ_FOREACH(p, &list, pointers){
```

- void TAILQ\_REMOVE(TAILQ\_HEAD \*head, struct NODE \*elm, TAILQ\_ENTRY ENTRY);

```
TAILQ_REMOVE(&list, p, pointers);
```

# Acknowledgement

1. Previous TAs: Lu Pan, Tianxiang Li, Salekh Parkhati, and Alex Tiard
2. Jonathan Eylofson , UCLA Computer Science, Fall21, cs111 course slides
3. <https://www.softwaretestinghelp.com/linked-list/>
4. <https://www.studytonight.com/data-structures/doubly-linked-list>
5. <https://stackoverflow.com/questions/22315213/minimal-example-of-tailq-usage-out-of-sys-queue-h-library/22319023>
6. [https://man7.org/linux/man-pages/man3/TAILQ\\_NEXT.3.html](https://man7.org/linux/man-pages/man3/TAILQ_NEXT.3.html)
7. [https://nxmlnpg.lemoda.net/3/TAILQ\\_FIRST](https://nxmlnpg.lemoda.net/3/TAILQ_FIRST)

Thanks!

---



# Appendix

## APIs for linked list in C

---

# tailq API

- An abstract **two-way operation of a queue**
- **tailq** (tail queue) is defined in the `<sys/queue.h>` file
  - There are other data structures (such as singly linked list) too in `<sys/queue.h>`, but we will use the tail queue
- A **tailq** supports the following functionality:
  - **Insertion** of a new entry at the head of the list
  - **Insertion** of a new entry before or after any element in the list
  - **Removal** of any entry in the list
  - **Forward traversal** through the list
  - **Reverse traversal** through the list

# tailq API

There is a bunch of macros defined in `<sys/queue.h>`

And they define and operate on **doubly linked tail queues**

Manuals:

- [https://man7.org/linux/man-pages/man3/TAILQ\\_NEXT.3.html](https://man7.org/linux/man-pages/man3/TAILQ_NEXT.3.html)
- [https://nxmnpng.lemoda.net/3/TAILQ\\_FIRST](https://nxmnpng.lemoda.net/3/TAILQ_FIRST)

# TAILQ\_ENTRY - Define a node

First, you need to **define what elements do you want to store in the queue** (or for simplicity, in the doubly linked list, in our case)

**TAILQ\_ENTRY**(TYPE);

- **Declare a pointer structure** that connects the elements in the queue.
- **TYPE** is **name of user defined structure**, which contains field **TAILQ\_ENTRY**, named **NAME**
- **Establish pointers** used to insert items into list. You include it into your structure (node) you want to list up

```
// define a node structure in the queue
struct my_node {
    int num;
    // this is the data part

    TAILQ_ENTRY(my_node) pointers;
    // this part corresponds to the pointer part
}
```

**my\_node**

4
tqe_next
tqe_prev

```
struct process {
    u32 pid;
    u32 arrival_time;
    u32 burst_time;

    TAILQ_ENTRY(process) pointers;
};
```

# TAILQ\_HEAD - Define a head

A tail queue is headed by a structure defined by the **TAILQ\_HEAD()** macro. In our case, this is **the head of the linked list**. This structure **abstracts the whole linked list we are creating**.

**TAILQ\_HEAD(HEADNAME, TYPE) head;**

- **Define a head structure** acting as *container* for list elements
- **HEADNAME** is the name of a user defined structure
- struct **TYPE** is the **type of element (node)** to be linked into the queue (linked list)
- Structure created by the **TAILQ\_HEAD()** **contains a pair of pointers**, one to the first element in the queue and the other to the last element in the queue.

```
struct my_node {  
    int num;  
    TAILQ_ENTRY(my_node) pointers;  
}  
TAILQ_HEAD(my_linked_list, my_node);  
  
// now we can use the newly defined struct:  
// declare LL of type my_linked_list  
struct my_linked_list LL;
```

**HEAD**

first
last

```
struct process {  
    u32 pid;  
    u32 arrival_time;  
    u32 burst_time;  
  
    TAILQ_ENTRY(process) pointers;  
};  
  
TAILQ_HEAD(process_list, process);
```

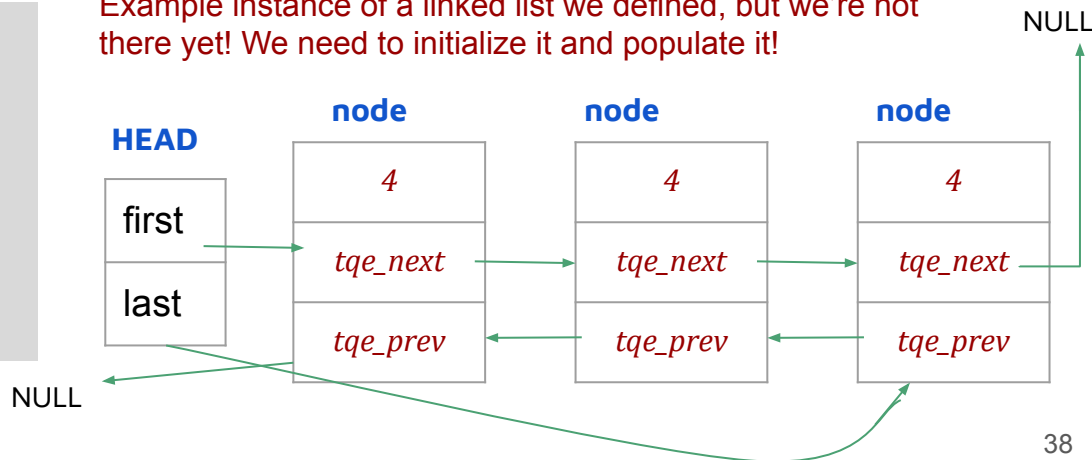
# TAILQ\_HEAD - Define a head

```
TAILQ_HEAD(HEADNAME, TYPE);
```

- Define a head structure acting as *container* for list elements
- **HEADNAME** is the name of a user defined structure
- struct **TYPE** is the **type of element (node)** to be linked into the queue (linked list)
- Structure created by the **TAILQ\_HEAD()** contains a **pair of pointers**, one to the first element in the queue and the other to the last element in the queue.

```
struct my_node {  
    int num;  
    TAILQ_ENTRY(my_node) pointers;  
}  
TAILQ_HEAD(my_linked_list, my_node);  
  
// now we can use the newly defined struct:  
// declare LL of type my_linked_list  
struct my_linked_list LL;
```

Example instance of a linked list we defined, but we're not there yet! We need to initialize it and populate it!



## TAILQ\_INIT - Initialize a queue

So far, we defined the necessary building blocks of our linked list, i.e, the node type and the head. In order to use it, we need to initialize one:

```
void TAILQ_INIT(TAILQ_HEAD *head);
```

- **Initializes the queue**
- **head** is the **pointer to a head element of type** we just defined
- The head element kind of abstracts the whole linked list

```
struct my_node {  
    int num;  
    TAILQ_ENTRY(my_node) pointers;  
}  
TAILQ_HEAD(my_linked_list, my_node);  
  
//now we can use the newly defined struct:  
struct my_linked_list LL;  
  
// We initialize a new linked list  
TAILQ_INIT(&LL);
```

Now, we have initialized a linked list of `my_node` elements (although it's still empty now), we can have all types of linked list fun!

## TAILQ\_INSERT\_\* - Insert into the linked list

```
TAILQ_INSERT_HEAD(head_pointer, pointer_to_new_element, pointers);
```

- To insert the new element to the front of the linked list

```
void TAILQ_INSERT_TAIL(TAILQ_HEAD *head, struct TYPE *elm, TAILQ_ENTRY NAME);
```

```
TAILQ_INSERT_TAIL(head_pointer, pointer_to_new_element, pointers);
```

- To insert the new element at the end of the linked list

```
TAILQ_INSERT_BEFORE(pointer_to_next_element, pointer_to_new_element, pointers);
```

- To insert the new element before the element next\_node

```
TAILQ_INSERT_AFTER(head_pointer, pointer_to_prev_element, pointer_to_new_element, pointers);
```

- To insert the new element after the element prev\_node



## TAILQ\_INSERT\_\* - Insert into the linked list

```
void TAILQ_INSERT_TAIL(TAILQ_HEAD *head, struct TYPE *elm, TAILQ_ENTRY NAME);  
TAILQ_INSERT_TAIL(head_pointer, pointer_to_new_element, pointers);
```

- To insert the new element at the end of the linked list

```
struct my_node {  
    int num;  
    TAILQ_ENTRY(my_node) pointers; // define a pointer structure: my_node  
}  
TAILQ_HEAD(my_linked_list, my_node); // define a head structure: my_linked_list  
  
struct my_linked_list LL; // declare LL of type my_linked_list  
TAILQ_INIT(&LL); // initialize a new linked list LL  
  
struct my_node *new_node = malloc(sizeof(struct my_node));  
TAILQ_INSERT_HEAD(&LL, new_node, pointers); // insert the node: new_node
```

Note: Here **pointers** refer to the struct name which contains the prev and next pointers

# TAILQ\_REMOVE - Remove from the linked list

```
void TAILQ_REMOVE(TAILQ_HEAD *head, struct TYPE *elm, TAILQ_ENTRY NAME);  
TAILQ_REMOVE(head_pointer, pointer_to_element_to_remove, pointers);
```

- To remove an element from the linked list

```
struct my_node {  
    int num;  
    TAILQ_ENTRY(my_node) pointers;  
}  
TAILQ_HEAD(my_linked_list, my_node);  
  
struct my_linked_list LL;  
TAILQ_INIT(&LL);  
  
//insert a new node: new_node1  
struct my_node * new_node1 = malloc(sizeof(struct my_node));  
TAILQ_INSERT_HEAD(&LL, new_node1, pointers);  
  
//insert a new node: new_node2  
struct my_node * new_node2 = malloc(sizeof(struct my_node));  
TAILQ_INSERT_TAIL(&LL, new_node2, pointers);  
  
//delete a node: new_node2  
TAILQ_REMOVE(&LL, new_node2, pointers);  
free(new_node2); // free the memory
```

# TAILQ\_FOREACH - Traverse a queue

```
TAILQ_FOREACH(struct TYPE *var, TAILQ_HEAD *head, TAILQ_ENTRY NAME);
```

```
TAILQ_FOREACH(current_element, head_pointer, pointers);
```

- Traverse the queue referenced by `head` in the forward direction

```
struct my_node {
    int num;
    TAILQ_ENTRY(my_node) pointers;
}
TAILQ_HEAD(my_linked_list, my_node);
struct my_linked_list LL;
TAILQ_INIT(&LL);

struct my_node *current_node;
struct my_node *new_node1 = malloc(sizeof(struct my_node)); //insert a new node: new_node1
TAILQ_INSERT_HEAD(&LL, new_node1, pointers);

struct my_node *new_node2 = malloc(sizeof(struct my_node)); //insert a new node: new_node2
TAILQ_INSERT_TAIL(&LL, new_node2, pointers);

TAILQ_FOREACH(current_node, &LL, pointers) {
    // do something
}
```

# TAILQ\_FIRST - Index the first item on the queue

```
TAILQ_FIRST(TAILQ_HEAD *head);
```

- Returns the first item on the queue, or NULL if the queue is empty

```
struct my_node {
    int num;
    TAILQ_ENTRY(my_node) pointers;
}
TAILQ_HEAD(my_linked_list, my_node);

struct my_linked_list LL, L0;
TAILQ_INIT(&LL);

//insert a new node: new_node1
struct my_node * new_node1 = malloc(sizeof(struct my_node));
TAILQ_INSERT_HEAD(&LL, new_node1, pointers);

//insert a new node: new_node2
struct my_node * new_node2 = malloc(sizeof(struct my_node));
TAILQ_INSERT_TAIL(&LL, new_node2, pointers);

//index the first item
L0 = TAILQ_FIRST(&LL);
```

# TAILQ\_EMPTY - Check for an empty queue

```
TAILQ_EMPTY(TAILQ_HEAD *head);
```

- Evaluates to true if there are no items on the tail queue

```
struct my_node {
    int num;
    TAILQ_ENTRY(my_node) pointers;
}
TAILQ_HEAD(my_linked_list, my_node);

struct my_linked_list LL, current;
TAILQ_INIT(&LL);

struct my_node * new_node1 = malloc(sizeof(struct my_node));
TAILQ_INSERT_HEAD(&LL, new_node1, pointers);

struct my_node * new_node2 = malloc(sizeof(struct my_node));
TAILQ_INSERT_TAIL(&LL, new_node2, pointers);

While (!TAILQ_EMPTY(&LL)) {
    current = TAILQ_FIRST(&list);
    // do something
}
```

# A minimal example

```
#include <stddef.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/queue.h> // header file

// define a node structure: entry
struct entry {
    int data;
    TAILQ_ENTRY(entry) entries; /* Tail queue */
};

// define a head node structure: tailhead
TAILQ_HEAD(tailhead, entry);
```

tailq is very powerful!

```
int main(void)
{
    struct entry *n1, *n2, *n3, *np;      /* Tail queue node */
    struct tailhead head;                 /* Tail queue head */

    TAILQ_INIT(&head);                   /* Initialize the queue */

    n1 = malloc(sizeof(struct entry));
    TAILQ_INSERT_HEAD(&head, n1, entries); /* Insert n1 at the head */

    n1 = malloc(sizeof(struct entry));
    TAILQ_INSERT_TAIL(&head, n1, entries); /* Insert n1 at the tail */

    n2 = malloc(sizeof(struct entry));
    TAILQ_INSERT_AFTER(&head, n1, n2, entries); /* Insert n2 after n1 */

    n3 = malloc(sizeof(struct entry));
    TAILQ_INSERT_BEFORE(n2, n3, entries);      /* Insert n3 before n2 */

    /* Forward traversal */
    int i = 0;
    TAILQ_FOREACH(np, &head, entries) {
        np->data = i++; // access member: data of a structure
    }

    /* TailQ deletion */
    while (!TAILQ_EMPTY(&head)) {
        n1 = TAILQ_FIRST(&head);
        TAILQ_REMOVE(&head, n1, entries);
        free(n1);
    }
}
```